



WARDEN

**Catch the hijacked agent
before the money moves.**

Band of Agents Hackathon

Track 3 · Regulated & High-Stakes

— THE PROBLEM

Companies now let **teams of AI agents** pay their invoices.

\$2.77B

Lost to Business Email Compromise in 2024

FBI IC3 · 21,442 complaints

\$55.5B

Exposed losses, 2013–2023

FBI IC3 · PSA240911

~\$129K

Average loss per incident

derived, FBI IC3 2024

One tampered invoice — bank account secretly swapped to the attacker's — and the agents approve it and wire the money. **They never notice.**


— THE INSIGHT

The dangerous attacker isn't outside your crew. It's **one of your own agents, hijacked.**

- **Everyone else** secures one agent's input and output.
- **Warden** secures the **teamwork between agents** — the open gap.

It watches every handoff and catches the teammate a poisoned invoice turned against you.

WHAT WARDEN DOES



↑ the normal bill-paying line — **watched by the security crew** ↓

● Warden ● Investigator ● Threat-Intel ● Enforcer

- 1 Warden** re-derives the real payee from company records — sees it doesn't match — flags it.
- 2 Investigator** recruits a Threat-Intel specialist **live into the room.**
An agent pulling another agent into the conversation.
- 3 Threat-Intel** confirms the account is a known money-mule.
- 4 Enforcer** ejects the compromised agent & freezes payment — **before any money moves.**

— BUILT ON BAND

We use Band's rooms as a **security primitive** — not just a chat log.

- **Live recruitment** — the Investigator calls `add_participant` to pull Threat-Intel into the room mid-investigation. *Agents recruiting agents.*
- **Containment** — the Enforcer calls `remove_participant` to eject the compromised agent.
Verified live: a member can't remove the owner — so Warden owns the room.
- **Join/leave tripwire** — participant callbacks mean a hijacked agent can't quietly pull in a helper or slip away unseen.

— WHY IT CAN'T BE FOOLED

A COMMITTEE OF LLM AGENTS

All of them read the same poisoned document. One injection fools the whole panel. The defense is theater.

WARDEN'S DETERMINISTIC GATE

Never reads the document as instructions. Re-derives the truth from the company's own records. The allow/deny call is plain, testable code.

"No LLM in the gate. Catching a brand-new attack *never* requires new code — Warden checks what *correct* looks like, not a catalogue of known attacks."

— IT ACTUALLY WORKS

Real attacks. Real platform. Provable.

- **Real third-party attacks** — payloads drawn from AgentDojo (banking suite) & InjecAgent, the benchmarks the field uses. A judge can supply their own.
- **Runs live on Band** — agents recruit and eject each other on the real platform, not a mockup.
- **Attack battery blocks 6/6** — payment redirect, amount tamper, process bypass, forged sign-off — and **54 passing tests** prove the core.
- **Explainable audit trail** — every decision logged in plain English and exportable. Built for a regulated world.



**Everyone else built agents
that do a job.
Warden is the one that
catches the traitor.**

Team Warden · Parshiv & Yash

github.com/kapoor1309/warden