



WARDEN

**Catch the hijacked agent
before the money moves.**

Band of Agents Hackathon

Track 3 · Regulated & High-Stakes

— THE SETUP

Today, AI agents pay company bills.

An agent is just a small AI worker. Companies now line up a few of them to handle invoices:

- One agent **reads** the invoice.
- One agent **checks** it against company records.
- One agent **approves** it.
- One agent **sends the money**.

This is fast and cheap. But it opens a door for thieves.

— THE PROBLEM

A thief sends a fake invoice.

- The invoice looks completely normal.
- But the bank account on it has been quietly changed to the **thief's account**.
- The agents do not notice. They approve it and send the money **straight to the thief**.

This is a real scam. Banks call it Business Email Compromise. In plain words: a fake invoice with a swapped bank account.

— WHY IT MATTERS

This is not small money.

\$2.77B

Stolen with fake invoices in 2024

FBI report (IC3), 21,442 cases

\$55.5B

Stolen over ten years (2013 to 2023)

FBI report (IC3)

~\$129K

Lost in a single typical attack

based on FBI 2024 numbers

Companies lose billions every year. One changed bank account is all it takes.

— THE GAP NOBODY COVERS

The real danger is **your own agent**.

- **Everyone else** protects one agent at a time: its input and its output.
- Nobody watches what happens **between** the agents.
- If a fake invoice tricks one of your agents, the whole team trusts it.
That tricked agent is now on the thief's side, and it looks just like a normal teammate.

— THE SOLUTION

Warden is a security team that watches the other agents.

It sits inside the same chat room as your agents and checks every step. It has four members:

● Warden, the guard

● Investigator, the detective

● Threat-Intel, the expert

● Enforcer, the bouncer

No money can move until Warden says it is safe.

— HOW IT WORKS, STEP BY STEP



The four agents above pay the bills. The security team watches them.

- 1** **Warden** checks the bank account on the invoice against the company's own records. It does not match, so Warden raises a warning.
- 2** **Investigator** needs a second opinion, so it brings a Threat-Intel expert into the room, live.
An agent calling in another agent for help.
- 3** **Threat-Intel** checks the account and confirms it is a known bad account used by criminals.
- 4** **Enforcer** removes the tricked agent and stops the payment. **No money leaves.**

— THE CLEVER PART

Why Warden cannot be fooled.

A TEAM OF AI AGENTS

They all read the same fake invoice. One trick fools all of them. They agree, and the money goes out.

WARDEN

It does not trust the invoice at all. It checks the company's own records instead. The yes or no answer is plain code, not an AI guess.

“Warden checks what a **correct payment** should look like. So it catches brand new tricks it has never seen, with no new code.”

— BUILT ON BAND

Our agents **really work together on Band.**

Band is the chat platform where the agents live. We use its rooms for real teamwork, not just messages:

- When an agent needs help, it **pulls another agent into the room.**
Agents bringing in agents.
- When an agent is caught, the Enforcer **removes it from the room.**
- If anyone joins or leaves, **Warden sees it at once.** A tricked agent cannot sneak in a helper or quietly slip away.

— PROOF IT WORKS

Real attacks. Real platform.

- We ran **6 real attacks** at it. Warden **blocked all 6**.
Payment redirect, amount changed, steps skipped, fake approval, and more.
- The attacks are **real, from public security tests** (AgentDojo and InjecAgent). A judge can try their own.
- It runs **live on the Band platform**, not just on a demo screen.
- **54 tests** check the core. Every decision is saved in a **signed report** that cannot be faked.



**Everyone else built agents
that do a job.**
**Warden is the one that
catches the thief.**

Team Warden · Parshiv & Yash

github.com/kapoor1309/warden